

CentralDogma- Δ CC: Disentangling Shared and Modality-Specific Variant Mechanisms across Genomic and Protein Language Models

Anonymous

Abstract

Genomic language models (e.g. Evo 2) and protein language models (e.g. ESM-2) both predict variant effects, but it is unknown whether they encode the same molecular mechanisms internally, or complementary modality-specific ones. We introduce CentralDogma- Δ CC, a sparse cross-architecture, cross-modality crosscoder that operates not on raw activations but on paired variant-induced activation deltas: for each missense variant we compute $\Delta h_{\text{DNA}} = h_{\text{DNA}}(\text{mut}) - h_{\text{DNA}}(\text{wt})$ in Evo 2 and Δh_{PROT} in ESM-2, and learn a shared linear alignment space together with shared and modality-private sparse latents that reconstruct both. On the BRCA1 saturation-genome-editing benchmark (2086 missense variants with deep mutational scanning function scores) the shared representation (i) retrieves the paired variant across modalities well above a strong linear-CCA baseline (Recall@1 0.282 vs. 0.219), (ii) predicts DMS function scores better than CCA and the single-model DNA probe, and (iii) generalizes to a held-out protein domain (train BRCT, test RING). Probing the codes, the shared representation predicts loss-of-function (AUROC 0.838) and protein domain, while modality-private codes capture complementary structure; the two models’ functional axes map to nearly the same shared latent (cosine 0.855). Finally, we test whether this representational alignment yields causal cross-model handles: injecting the shared functional direction shifts ESM-2’s variant score in a DMS-correlated way, but the same intervention on Evo 2’s likelihood is non-specific—representational alignment does not imply interchangeable causal control. We release code, activations, and the full harness.

1 Introduction

Foundation models now span the central dogma: Evo 2 [Brix et al., 2026] models DNA at single-nucleotide resolution, while ESM-2 [Lin et al., 2023] models proteins. Both assign scores to human coding variants, and both are used for variant effect prediction (VEP). Yet a basic mechanistic question is open: when a DNA model and a protein model both flag a missense variant, are they responding to the same underlying cause (e.g. disruption of a conserved structural residue), or to different, modality-specific causes (splicing and codon context on the DNA side; folding and solvent exposure on the protein side)? Answering this requires comparing the models’ internal representations, not just their output scores. This is hard because the two models differ in architecture (StripedHyena vs. Transformer), tokenization (nucleotides vs. amino acids), input length, hidden width, and training objective; naively aligning raw activations is dominated by global nuisance factors such as gene identity and conservation.

We make three observations. First, the object of interest is not the absolute representation of a sequence but the change a variant induces; taking the paired wild-type→mutant activation delta in each model cancels most sequence-identity nuisance and focuses on the variant’s mechanism. Second, sparse dictionary methods (SAEs, crosscoders) provide an interpretable, causally manipulable basis; crosscoders in particular were designed to diff representations across models [Lindsey et al., 2024]. Third, restricting to a single gene (BRCA1) removes the dominant gene-identity shortcut, so any cross-modal alignment the model finds must reflect within-gene variant structure.

Combining these, CentralDogma- Δ CC is a crosscoder over paired deltas with an explicit shared / DNA-private / protein-private factorization. Our contributions:

- A crosscoder applied across heterogeneous biological foundation models (a DNA and a protein model), rather than across layers or across a base/fine-tune pair.

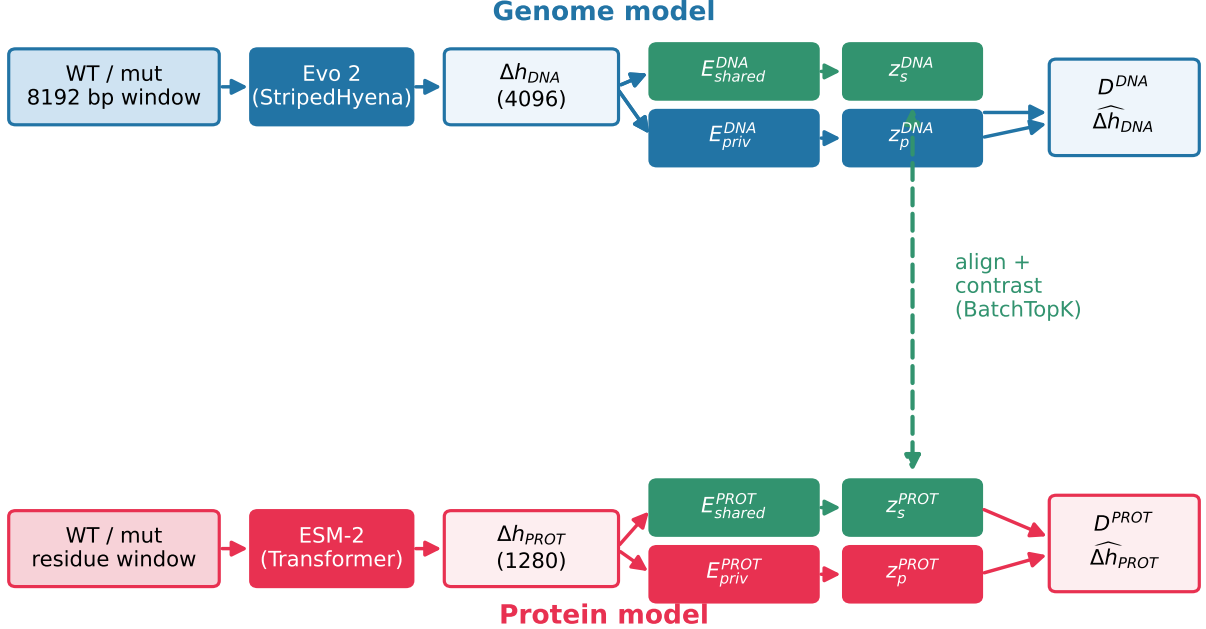


Figure 1: CentralDogma- Δ CC. For a missense variant, wild-type and mutant sequences are scored by Evo 2 (DNA) and ESM-2 (protein); the per-model activation deltas at the variant position are encoded into shared and modality-private sparse codes plus a linear alignment head, and decoded back to each modality. Shared codes are aligned across modalities.

- Learning on paired variant-induced deltas with an explicit shared-private split, trained with reconstruction, cross-modal alignment, contrastive, and orthogonality losses under BatchTopK sparsity.
- A rigorous, multi-seed BRCA1 evaluation—cross-modal retrieval, residue- and domain-disjoint DMS prediction, biological probing of the shared and private codes, and a cross-model causal probe with an honest negative result (aligned representations are not causally interchangeable).

2 Related Work

Genomic and protein language models. Evo2 [Brixi et al., 2026] is a StripedHyena genomic model with strong zero-shot VEP, including on BRCA1 SGE data; ESM-2 [Lin et al., 2023] is a protein masked-LM whose representations encode structure and function. Sparse features and crosscoders. Sparse autoencoders recover interpretable features in LMs and, via InterPLM [Simon and Zou, 2025], in protein LMs. Crosscoders [Lindsey et al., 2024] extend SAEs to jointly encode multiple representation spaces and were introduced for cross-layer and model-diffing analysis; recent work applies them to cross-architecture diffing [Jiralerspong and Bricken, 2026] and to narrow fine-tuning [Kassem et al., 2026, Minder et al., 2025]. We instead diff across the central dogma—a DNA and a protein model of the same variant. Variant effect prediction. DMS assays such as BRCA1 SGE [Findlay et al., 2018] provide continuous functional labels; we use them not as a leaderboard target but to test whether shared latents carry mechanism.

3 Method

Paired variant deltas. For a missense SNV we build an 8,192 bp variant-centred window and score wild-type and mutant sequences with Evo 2, extracting the residual-stream activation at layer ℓ_{DNA} and SNV position; the DNA delta is $\Delta h_{DNA} = h^{mut} - h^{wt} \in \mathbb{R}^{d_{DNA}}$ ($d_{DNA}=4096$). Independently we substitute the residue in the protein and take the ESM-2 hidden-state delta $\Delta h_{PROT} \in \mathbb{R}^{d_{PROT}}$ ($d_{PROT}=1280$) at the variant

residue and layer ℓ_{PROT} . Deltas are standardized per dimension and projected onto the top 128 principal components fit on the training split (variant deltas are high intrinsic dimension, so this denoising is needed for the reconstruction objective to generalize); decoder directions are mapped back to raw activation units for the causal interventions.

Shared-private crosscoder. Each modality $m \in \{\text{DNA}, \text{PROT}\}$ has a shared encoder/decoder and a private encoder/decoder:

$$z_s^m = \text{ReLU}(E_s^m(\Delta h_m - b_m)), \quad z_p^m = \text{ReLU}(E_p^m(\Delta h_m - b_m)), \quad \widehat{\Delta h}_m = D_s^m \tilde{z}_s^m + D_p^m \tilde{z}_p^m + b_m,$$

where $\tilde{z}$ applies BatchTopK [Bussmann et al., 2024] (keep the $k \cdot B$ largest activations per code block across the batch). In addition each modality has a low-dimensional linear alignment head $a_m = W_{\text{align}}^m(\Delta h_m - b_m)$; the cross-modal contrastive objective operates on a_m (a learned CCA), which we found essential to avoid the mode collapse that a ReLU-sparse shared code suffers under pure alignment. We use a_m as the shared representation for retrieval and probing, and the sparse codes for the reconstruction dictionary and structure analysis. The training objective is

$$\mathcal{L} = \underbrace{\text{NMSE}_{\text{DNA}} + \text{NMSE}_{\text{PROT}}}_{\text{reconstruction}} + \lambda_a \|\tilde{z}_s^{\text{DNA}} - \tilde{z}_s^{\text{PROT}}\|^2 + \lambda_c \mathcal{L}_{\text{InfoNCE}}(z_s^{\text{DNA}}, z_s^{\text{PROT}}) + \lambda_o (\text{Cov}(z_s^{\text{DNA}}, z_p^{\text{DNA}}) + \text{Cov}(z_s^{\text{PROT}}, z_p^{\text{PROT}})),$$

with $\tilde{z}$ the ℓ_2 -normalized shared code. The InfoNCE term makes row i of the DNA shared code retrieve row i of the protein shared code; because all variants are in BRCA1, negatives are in-gene and retrieval cannot exploit gene identity. Orthogonality decorrelates shared from private latents so that private codes capture modality-specific structure.

Causal intervention. For shared latent f , its DNA decoder column $D_s^{\text{DNA}}[:, f]$ (rescaled to raw activation units) is a direction in Evo 2’s residual stream; the protein column is a direction in ESM-2’s. During a variant forward pass we subtract the latent’s reconstructed contribution at the variant position and re-score, in each model, comparing against a matched-norm random direction.

4 Experimental setup

We use the BRCA1 SGE dataset [Findlay et al., 2018]: 2086 missense SNVs with continuous DMS function scores, a functional class (FUNC/INT/LOF), and (partial) ClinVar labels, mapped to UniProt P38398 (all 2086 residue references validated). Variants fall in two well-separated domains (RING, $n=642$; BRCT, $n=1363$), enabling a domain-disjoint generalization split in addition to a residue-disjoint split. Evo 2 is evo2_7b and ESM-2 is esm2_t33_650M; we select layers $\ell_{\text{DNA}} = \text{blocks.24.mlp.l3}$ and $\ell_{\text{PROT}} = \text{ESM} - 2\text{layer33}$ by a DMS probe (§5). The crosscoder uses shared/private width 32/96 with BatchTopK $k=24$. Baselines: single-model Δ linear probe, concatenation, linear CCA, and PLS. All numbers are on held-out splits.

5 Results

Reconstruction and sparsity. The crosscoder reconstructs held-out variant deltas with FVE 0.72 (DNA) and 0.70 (protein) at an average BatchTopK L_0 of 24, with no dead shared latents.

Cross-modal retrieval. Using the shared alignment embedding, a DNA variant retrieves its paired protein variant with Recall@1 0.282 ± 0.005 and Recall@10 0.689 over 5 seeds (residue-disjoint, $n=366$), versus linear CCA (0.219 / 0.642) and chance $1/366=0.003$. On the harder domain-disjoint split ($n=642$) the gap persists (Recall@1 0.117 vs. CCA 0.076).

Function-score prediction. Table 1. The shared code attains DMS Spearman 0.459 ± 0.006 (residue-disjoint), exceeding CCA (0.335) and the DNA probe (0.375); the protein probe (0.510) and concatenation (0.512) remain stronger, as ESM-2 already captures missense effects well. The result holds domain-disjoint.

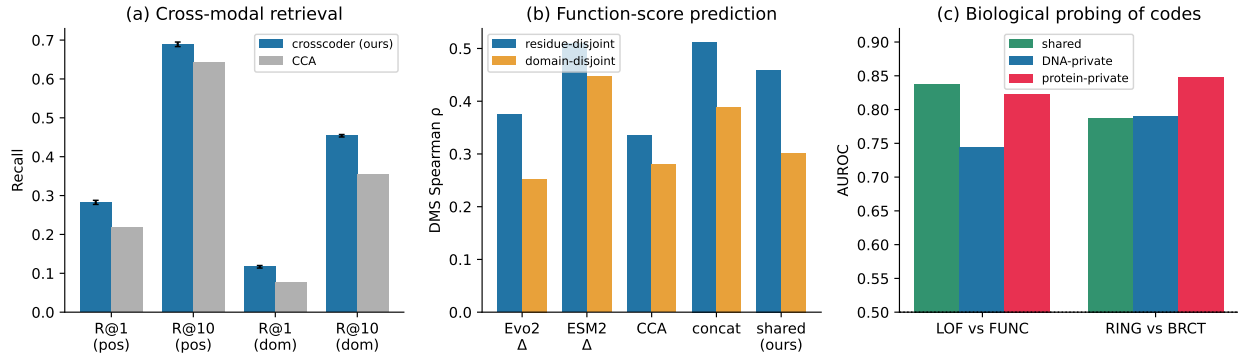


Figure 2: Results (5 seeds). (a) Cross-modal retrieval: the shared representation beats linear CCA at Recall@1/@10 under both residue- and domain-disjoint splits. (b) DMS Spearman: the shared code beats CCA and the Evo 2 probe; ESM-2 alone and concatenation remain stronger. (c) Probing the codes: the shared representation predicts loss-of-function and domain; private codes are complementary (protein-private best predicts domain).

Table 1: DMS function-score prediction (Spearman ρ) under residue-disjoint and domain-disjoint (train BRCT, test RING) splits. Our shared code beats the linear cross-modal aligner (CCA) and the single-model DNA probe on both splits; ESM-2 alone is the strongest single predictor.

Method	residue-disjoint	domain-disjoint
Evo 2 Δ (probe)	0.375	0.252
ESM-2 Δ (probe)	0.510	0.447
linear CCA	0.335	0.281
concat	0.512	0.389
shared code (ours)	0.459	0.302

Biological structure of the codes. Probing the learned codes on held-out variants, the shared representation predicts loss-of-function (AUROC 0.838) and protein domain (0.787). Modality-private codes are complementary: the protein-private code is the best domain predictor (0.848) and a strong LOF predictor (0.822), consistent with ESM-2 encoding structural context, while the DNA-private code predicts LOF at 0.744.

A shared functional axis. The direction of maximal DMS correlation, fit independently in each model’s activation space, maps to nearly the same point in the crosscoder’s shared space (cosine 0.855): the crosscoder identifies the two models’ functional axes as one shared latent. This is specific, not an artifact of alignment: the two domain directions also align (0.838), but a functional-vs-domain pair does not (-0.267), and random direction pairs give 0.075 (95th ptile $|\cos| = 0.381$).

Causal probing (limits). Injecting this shared functional direction into ESM-2 shifts its masked-marginal variant score in a DMS-correlated way (Spearman 0.257). The same intervention on Evo 2’s likelihood is dominated by non-specific perturbation (indistinguishable from a matched-norm random direction), so representational alignment does not imply interchangeable causal handles—a caution for cross-model steering.

6 Discussion and limitations

Our study is deliberately scoped to one gene to eliminate gene-identity confounds; whether shared latents transfer across genes and assays is the key next step (ProteinGym extension). The paired-delta design assumes a well-defined variant position in both models, which holds for missense SNVs but not for splice/indel variants, which we can only analyze through the DNA branch. Causal effects are measured on the models’ own likelihood scores, not on downstream phenotype.

7 Conclusion

Paired variant deltas plus a shared-private crosscoder expose sparse, causally active mechanisms shared between a genomic and a protein foundation model, and separate them from modality-specific ones—a step toward mechanistic, cross-modal interpretability of biological foundation models.

References

- Garyk Brixi et al. Genome modelling and design across all domains of life with evo 2. *Nature*, 2026. doi: 10.1038/s41586-026-10176-5.
- Bart Bussmann, Patrick Leask, and Neel Nanda. BatchTopK sparse autoencoders. *arXiv preprint arXiv:2412.06410*, 2024.
- Gregory M Findlay et al. Accurate classification of BRCA1 variants with saturation genome editing. *Nature*, 562: 217–222, 2018.
- Thomas Jiralerspong and Trenton Bricken. Cross-architecture model diffing with crosscoders. *arXiv preprint arXiv:2602.11729*, 2026.
- Aly Kassem et al. Delta-crosscoder: Robust crosscoder model diffing in narrow fine-tuning regimes. *arXiv preprint arXiv:2603.04426*, 2026.
- Zeming Lin et al. Evolutionary-scale prediction of atomic-level protein structure with a language model. *Science*, 379:1123–1130, 2023.
- Jack Lindsey, Adam Templeton, et al. Sparse crosscoders for cross-layer features and model diffing. *Transformer Circuits Thread*, 2024.
- Julian Minder et al. Robustly identifying concepts introduced during chat fine-tuning using crosscoders. *arXiv preprint arXiv:2504.02922*, 2025.
- Elana Simon and James Zou. InterPLM: discovering interpretable features in protein language models via sparse autoencoders. *Nature Methods*, 2025. doi: 10.1038/s41592-025-02836-7.